

Inference cost is ~ 40 ms, near-constant in q

